

CoArena Open-Source Computer-Use Assistant: Product, Data and Gym Strategy

Executive summary

Recommendation: proceed, but do not build the proposition as “a free assistant that collects users’ computers”. The strategically strong version is a **free, open-source, model-agnostic computer-use client whose default modes are private, with a separately consented trajectory-contribution programme that converts selected real workflows into synthetic, reproducible CoArena Gym environments.** The distinction is fundamental. Raw desktop trajectories are unusually valuable for understanding the real distribution of computer-use tasks, but they are also unusually dangerous: screenshots and typed actions can contain credentials, financial information, health data, private messages, confidential work material and third-party personal information. California explicitly treats account credentials, financial information and the contents of mail, email and text messages as sensitive personal information; GDPR principles require purpose limitation, minimisation, storage limitation and privacy by design. ¹

There is strong strategic evidence for the underlying thesis. Current frontier computer-use agents remain far from reliable on realistic long-horizon work: OSWorld 2.0 contains 108 workflows whose human median completion time is about 1.6 hours, and its best reported system completed only 20.6% under the binary 500-step metric. Failures were less about elementary clicking than about losing constraints, missing information arriving mid-task, guessing rather than asking, and failing to verify.

² WeaveBench similarly finds that real workflows benefit from trajectory-level evidence: its 114 tasks use real-world artifacts, frozen environments and trajectory-aware judging; the best reported model/runtime combination reaches 41.2% PassRate, and GUI-only or CLI-only ablations stay at or below 3.5% for its deliberately hybrid tasks. ³ ComputerRL, meanwhile, shows why reproducible executable environments matter for training: it builds large-scale desktop environments and trains with end-to-end reinforcement learning and verifiable rewards. ⁴

That means the highest-value asset is probably **not the raw trajectory corpus**. It is the transformation:

**real user intent → observed workflow distribution → sanitised task abstraction →
synthetic fixtures → reproducible environment → deterministic/trajectory-aware
grader → Gym task → model evaluation/RL**

That creates a far more defensible company than selling screenshots. The assistant becomes CoArena’s **real-world task-distribution sensor**; CoArena becomes the benchmark and comparative intelligence layer; Gym becomes the monetisable training/evaluation product.

This also fits CoArena’s existing work unusually well. Internal materials indicate that Coasty has already developed an **MIT-licensed “Open Cowork” desktop agent** that operates through screenshots, mouse and keyboard, separates the agent loop from any single model, supports local Ollama plus hosted providers including OpenAI, Anthropic, Google, xAI, Mistral and OpenRouter, and lets users watch and interrupt execution. CoArena also already has Gym/RL-environment work in progress, so this should be

treated as a productisation, privacy-hardening and data-pipeline problem rather than a greenfield CUA research project.

There is also direct internal evidence that long-horizon trajectory data has commercial demand. In the Conifer work, the requested starting dataset was 1,000 pairs of long-horizon, non-verifiable tasks outside conventional maths/coding benchmarks, with model, query, response, token/cost information, user preference, tool calls and retries. At the same time, CoArena's own privacy requirements to Conifer explicitly asked for **no training on prompts, completions, screenshots or tool calls, no content sharing beyond the serving provider, metadata-only logging, and deletion on request/termination**. That tension is instructive: buyers want rich trajectories, but sophisticated users and partners increasingly expect strict content boundaries. The product must reconcile those expectations through **clear separation of private inference from voluntary data contribution**.

My assessment is therefore:

Dimension	Assessment	Implication
Strategic fit with CoArena	Very high	Directly supplies task distribution and Gym candidates.
Consumer utility today	Moderate	Excellent for bounded tasks; long-horizon reliability remains poor. ⁵
Data value	Very high	Intent, state, actions, corrections and outcomes are richer than preference votes alone.
Raw-data monetisation attractiveness	Low	Privacy, contractual, third-party-data and reputational liability swamp much of the upside. ⁶
Environment/Gym monetisation attractiveness	Very high	Converts private, messy observations into repeatable training and evaluation infrastructure. ⁷
Privacy/security risk	Extremely high	Accessibility + screen-recording permissions effectively expose the user's digital life.
MVP feasibility	High	Core agent harness already appears to exist internally.
"Free frontier inference for everyone" economics	Poor	Long action loops repeatedly send image/context tokens; BYOM/local should carry most free usage. ⁸

The decisive product principle should be: *private by default, contribution by choice, environments by design*.

The report assumes **macOS, Windows and Linux are all ultimately in scope and there is no platform constraint specified by the user**. Nevertheless, I recommend shipping the first production-quality beta on one primary OS—probably macOS given the existing power-user/developer wedge—then adding Windows and Linux. Supporting all three simultaneously in the first twelve weeks would multiply permission, capture, input, DPI/multi-monitor and packaging QA without increasing the validity of the initial product/data thesis.

Product thesis, users and competitive landscape

The product should be positioned less like “ChatGPT with access to your screen” and more like **OpenCode for the rest of your computer**:

An open-source assistant that can use software the same way you do—by seeing the screen, clicking, scrolling and typing. Bring any model, run locally where possible, inspect every action, stop it instantly, and choose exactly which runs you contribute to computer-use research.

OpenCode is an unusually good product-design precedent because it separates a useful open-source client from model choice. It supports terminal, IDE and desktop interfaces, advertises more than 75 providers including local models, and states that it does not store users’ code or context; its enterprise documentation says processing happens locally or via direct calls to the chosen AI provider except for deliberately enabled sharing. ⁹ The CoArena analogue should preserve that trust model while adding GUI actuation.

The “GUI-only” restriction is strategically useful but should be understood as a **research/product mode, not a religious architecture constraint**. A pure screenshot → action loop maximises signal about the exact capability CoArena wants to measure. Every successful or failed task becomes evidence about visual grounding, cursor control, planning, recovery and verification. But WeaveBench provides a warning: for some real workflows, GUI and structured/programmatic interfaces are complementary; its deliberately hybrid tasks collapse to $\leq 3.5\%$ success when restricted to either channel alone. ³ Thus I would expose **“GUI Research Mode”** as the canonical CoArena trajectory mode and preserve pure-GUI trajectories there, while leaving room for a later “Fast/Hybrid Mode” that can use structured tools when consumer utility matters more than clean CUA measurement.

Primary personas. The best first users are not mainstream consumers trusting an agent with their bank account. They are technically sophisticated people who understand model imperfections and already tolerate experimental agents.

Persona	Jobs to be done	Why they are attractive	Early risk
AI power user / early adopter	“Do this repetitive thing across Chrome, Finder, Sheets and Notes.”	High tolerance for errors; naturally shares demos and feedback.	May optimise for novelty rather than retained utility.
Developer / agent researcher	Compare Claude/Gemini/GPT/local models on the same desktop task.	Directly aligned with CoArena; valuable provider/model telemetry.	Can bias dataset toward technical workflows.
Knowledge worker with repetitive low-risk work	File organisation, data transfer, browser research, document formatting.	Real economically useful task distribution.	Confidential work content frequently appears on screen.
Privacy/local-AI enthusiast	Run an assistant through Ollama or self-hosted inference without a central service.	Strong open-source constituency; validates private mode.	Local CUA model quality/hardware can be limiting.

Persona	Jobs to be done	Why they are attractive	Early risk
Eval/model lab researcher	Run repeatable task packs and compare agents.	Natural bridge into paid Gym/evals.	Different UX from consumer assistant.
Voluntary research contributor	Donate specific runs for credits or research impact.	Produces labelled real-world task data.	Incentives can distort the natural task distribution.

Explicit non-targets for the first beta should be regulated workflows, unattended financial transactions, healthcare record handling, employer surveillance, privileged enterprise desktops, minors and users expecting perfect autonomous delegation. Google's current computer-use guidance itself treats account logins, sensitive health/financial/government records, personal identifiers, file/data transfer and consequential actions as requiring special handling or user confirmation, and recommends secure execution environments, guardrails, allow/block lists and detailed auditability. ¹⁰ OSWorld 2.0's reliability results make fully unattended positioning commercially dangerous today. ²

Competitive landscape. No single current product combines all of the desired attributes—general desktop GUI actuation, fully open client, broad provider choice, local/private operation, a transparent voluntary contribution programme, and a downstream environmentisation/Gym business. That is the gap.

Existing solution	General GUI control	Open-source	Privacy model	Model providers	Data collection posture	Maturity / CoArena gap
OpenCode	No general desktop GUI; coding in terminal/IDE/desktop client	Yes	Says code/context are not stored; local or direct provider requests; optional share is separate	75+ providers, local models, existing subscriptions	Privacy-first; sharing is explicit	Very mature product pattern to copy, but focused on code rather than screen actuation. ⁹
Raycast AI	Partial observation , not a general pixel-driven mouse/keyboard agent; Screen Awareness captures focused-window screenshot/accessibility context	No for core app	Capture occurs when invoked; Screen Awareness says it does not background-capture or retain the capture; AI privacy controls include local/custom/BYOK options	Broad hosted models, local models, custom providers and BYOK/BYOS	Operational metadata and explicitly submitted feedback are distinguished from ordinary interactions	Very polished UX/trust benchmark; weaker as a pure CUA trajectory source. ¹¹

Existing solution	General GUI control	Open-source	Privacy model	Model providers	Data collection posture	Maturity / CoArena gap
MacGPT	Partial: global UI and inline typing, not autonomous screenshot/click loop	No	App connects directly to OpenAI in API mode; its published privacy posture says it does not store/track usage	Primarily OpenAI/ChatGPT	Not designed as a trajectory collector	Strong simple launcher UX but narrow, and parts of the public feature/FAQ copy are visibly dated. ¹²
AutoGPT	No general pixel desktop loop as core product; mainly workflows/integrations	Mixed: Classic MIT; current platform core uses Polyform Shield/ source-available terms	Self-hosted version runs on user infrastructure; hosted service manages credentials/ infrastructure	Hundreds of models; 45+ connected platforms claimed	Depends on hosted versus self-hosted deployment	Mature agent/ workflow ecosystem, but API/ integration-centric rather than GUI-native. ¹³
UI-TARS Desktop	Yes: screenshot/ vision plus mouse and keyboard on local/remote computer/ browser	Yes, Apache-2.0	Repository describes fully local processing for local mode; remote operators also exist	Desktop is centred on UI-TARS/ Seed vision models; broader Agent TARS supports additional providers	No CoArena-style user trajectory donation/ marketplace is the core proposition	Closest technical competitor; validates open-source local GUI agent demand, but lacks CoArena's proposed data→eval→Gym loop. ¹⁴
Browser Use	Browser only, rather than full desktop; local or cloud browser control	Yes, MIT for library	Can be run locally/self-hosted; hosted browser/ agent services have separate recordings/ data policies	Direct OpenAI/ Anthropic/ Google wrappers, Browser Use gateway, Ollama/local and others	Local library can avoid their cloud; hosted service includes managed recordings/ data-policy infrastructure	Strong developer maturity and huge browser wedge; gap is native desktop breadth plus CoArena environmentisation. ¹⁵

Two adjacent systems matter even though they are not direct consumer competitors.

WeaveBench is a blueprint for the downstream asset. Its tasks are sourced from real user requests and public artifacts, packaged with initial environments, seed data, required assets, expected deliverables, reference trajectories and verification anchors. Runs occur in a frozen Linux VM that resets after completion, and grading inspects artifacts, screenshots, logs and trajectory evidence rather than trusting a final answer. ³ CoArena should borrow this *structure*, then make task sourcing more systematically tied to actual personal-assistant usage.

Conifer is better viewed as infrastructure/partner/buyer signal than a CUA competitor. Internally it has asked CoArena for long-horizon, non-verifiable data including query/response, model, costs, user selection, tool calls and retries, while the CoArena side simultaneously requested strict no-training/no-content-storage treatment for gateway traffic. This supports a critical conclusion: **the market values trajectory metadata, but trust is likely to depend on separating inference traffic from data licensing.**

The defensible gap, therefore, is not “we also have a desktop agent”. UI-TARS already demonstrates that. The gap is:

OpenCode-like distribution and model choice + Raycast-like trust UX + UI-TARS-like GUI control + CoArena comparative evaluation + WeaveBench-style reproducibility + a consented real-world task-distribution engine.

MVP product and user experience

The MVP should feel tiny. A user installs it, grants clearly explained permissions, picks where inference runs, presses a global shortcut, states an objective, watches the cursor move, and can stop it at any point.

The client should never pretend that accessibility or screen-recording permission is innocuous. On Windows, the native capture API deliberately presents system UI for users to select what is captured and marks actively captured content; on Linux/Wayland, the XDG Remote Desktop portal has explicit session/device selection and a user-facing start dialog. ¹⁶ On macOS, ScreenCaptureKit is the native capture framework, while input automation requires the corresponding accessibility/event permissions. ¹⁷ These permissions are effectively the keys to the user’s desktop, so CoArena’s UI should make them continuously visible, reversible and understandable.

Recommended first-run flow

Install → explain capabilities → choose privacy/inference mode → choose model/provider → grant screen permission → grant input permission → configure protected apps → run safe tutorial → enable hotkey

The safe tutorial should operate on a CoArena-owned local demo window, not the user’s real mail or files. For example: “Move the red card into the Completed column and type ‘done’.” That proves capture, coordinate mapping, clicking, dragging, typing, kill-switch handling and model integration without exposing meaningful information.

The most important screen in onboarding is not API-key entry. It is the **mode selector**:

Mode	What leaves the computer	Who pays for inference	CoArena trajectory collection	Recommended default
Local Private	Nothing except optional non-content crash telemetry	User hardware	None	Best privacy; default when supported
BYOM Private	Screenshots/prompts go directly to the chosen model provider	User/API subscription	No content; optional aggregate operational metadata	Default cloud mode
CoArena Sponsored	Screenshots/prompts go through the disclosed inference path	CoArena	Only a run explicitly enrolled in the research programme is donation-eligible	Optional
Donate this Run	A reviewed/sanitised trajectory is uploaded after explicit confirmation	Either	Yes, exactly the selected run	Best initial data-acquisition mechanism
Auto-Contribute Research Mode	Eligible runs are sanitised/uploaded under a persistent opt-in, with visible per-run notice and immediate revocation	Either	Yes, under versioned consent	P1, not launch default

This architecture deliberately avoids making trajectory collection a prerequisite for the core free client. That is not merely a branding preference. The European Commission's GDPR guidance states that valid consent must be freely given, informed, specific and affirmative, that withdrawal must be as easy as giving consent, and that consent may not be freely given where unnecessary data processing is made a precondition of the service. ¹⁸ A free local/BYOM mode therefore materially strengthens both trust and the legal argument that research contribution is genuinely optional.

The run UX should be:

Hotkey → task → scope preview → Start → live screenshot/action stream → confirmation gates as needed → Complete/Failed → user rating/correction → optional donation review

The **scope preview** need not be a slow model-generated plan. It should say, in deterministic language:

"CoArena will be able to see the selected display and control your mouse and keyboard. It will pause before sending messages, submitting forms, making purchases, deleting data, changing account/security settings, or transmitting detected sensitive information."

During execution, a persistent floating indicator should show **Seeing, Thinking, Clicking/Typing, Waiting for you, or Paused**. A global stop key must work even if the CoArena window is unfocused. “Pause” and “Stop” must be implemented outside the model loop so a stuck API request cannot disable them.

The action vocabulary for GUI Research Mode should stay deliberately small:

capture, click, double_click, right_click, move, drag, scroll, type_text, key, hotkey, wait, request_user, done.

This is close to the minimal primitive approach used by WeaveBench, whose GUI plugin exposes screenshot plus a small set of pyautogui-backed actions. ³ Anthropic’s current computer-use tooling similarly centres on screenshots and explicit interaction primitives, and its documentation places execution in a client-controlled environment rather than giving the model direct operating-system access. ¹⁹

For the **MVP, no shell, file API, browser DOM, email API or hidden accessibility tree should be exposed to the decision-making model in Research Mode**. That keeps trajectories genuinely interpretable as computer-use performance. An accessibility tree can still be used **locally by the safety subsystem** to recognise password fields or sensitive controls, provided it is explicitly excluded from the model observation and tagged separately in the trace. That provides a useful compromise between scientific cleanliness and user safety.

Permissions should be narrower than OS permissions. Even when the operating system forces CoArena to receive broad capture/accessibility privileges, the application can implement its own exclusion layer:

User control	MVP behaviour
Protected apps	Password managers, banking apps and selected enterprise apps are blocked by default or user-configurable.
Protected sites	Domain blocklist; financial, identity, health and admin domains can default to takeover/confirmation.
Display scope	User selects the monitor/window where the run may operate where OS APIs permit.
Sensitive-input mode	Before password/MFA/secret entry, agent pauses, screenshots stop being persisted, and user takes control.
External side effects	“Send”, “Post”, “Buy”, “Pay”, “Delete”, “Submit”, security-setting changes and data sharing require confirmation.
Clipboard	Off for the agent by default; explicit opt-in if later added.
Maximum run	Hard turn/action/time budgets.
Global kill switch	Always available independent of provider/model process.

Google’s own computer-use example embodies the same “prepare everything, then stop before the consequential click” design for payments and posts, and recommends secure execution, allow/block

lists, sanitisation, guardrails and logging. ¹⁰ This should be a hard client policy, not merely a sentence in the system prompt.

The donation experience is the centrepiece of the data strategy. After a run, show:

Help improve computer-use models?

CoArena found 34 screenshots and 29 actions in this run.

We detected 3 potentially sensitive regions and will remove them before upload.

You can review the run before contributing it.

The user should be offered three granular options:

Contribution level	Data
Task statistics only	Model/provider, OS, application classes, duration, action counts, error categories, cost/tokens, user-reported success.
Sanitised trajectory	Task instruction, selected/redacted frames, executable actions, intervention/correction events, final outcome.
Temporary research raw run	Short-retention encrypted raw trajectory for debugging/sanitiser research; separate explicit consent, not needed for routine donation.

A timeline review should let the user scrub frames, remove steps or cancel the donation entirely.

“Donate” should never be pre-ticked.

The **core MVP feature set** should therefore be narrower than a consumer “super assistant”:

Priority	Capability
P0	Global command bar; screenshot→action loop; cursor/keyboard control; live run viewer; pause/stop; task completion.
P0	macOS production-quality permissions/capture/input; Windows technical path behind feature flag.
P0	Anthropic, Google, OpenAI-compatible and generic OpenRouter-style provider adapters; local Ollama adapter.
P0	Private local/BYOM mode with zero CoArena content upload.
P0	Encrypted local run recorder with per-run delete.
P0	Side-effect policy engine and confirmation gates.
P0	Post-run optional donation flow with frame review.
P0	Minimal content-free telemetry and cost/token accounting.
P1	Local secret/PII detection and visual redaction.
P1	CoArena encrypted contribution ingest and deletion endpoint.
P1	Windows production support, multi-monitor/DPI hardening.


Priority	Capability
P1	Task clustering/environmentisation console.
Later	Persistent autonomous workflows, scheduling, background operation, hybrid API mode, enterprise fleet controls.

Telemetry should be intentionally boring. In normal private mode, the default CoArena event should contain app version, OS/version, provider/model ID, run duration, model latency, number of screenshots/actions, retry/error codes, token counts where providers expose them, safety-gate counts, user-reported success and a random pseudonymous installation identifier. It should **not** contain prompt text, screenshots, typed strings, URLs, window titles, clipboard contents, filenames or model reasoning. This is also directionally consistent with the privacy standard CoArena itself requested from Conifer: metadata acceptable, content not.

Raycast is a particularly useful UX benchmark here: its Screen Awareness is explicitly invoked, tells users what was captured, works with explicit OS permissions and states that it does not continuously capture the screen in the background. ¹¹ CoArena should strive for the same feeling of **observable**, **intentional access**, even though its actual execution loop is more powerful.

Technical architecture and Gym data engine

The simplest robust architecture is a **local-first desktop harness with provider adapters and a completely separate contribution pipeline**. The inference path should not require CoArena's data servers.



The architecture should preserve a very strong invariant:

The client can function fully without the CoArena trajectory-ingest service being reachable.

That gives the open-source codebase an independently useful reason to exist and makes network inspection auditable. OpenCode's current enterprise architecture uses the same broad trust idea: local processing or direct provider requests, with OpenCode itself not storing context unless users deliberately share. ²⁰

Client stack. Because an internal Open Cowork prototype already exists, the first instruction should be **do not rewrite the UI framework simply to fit a preferred architecture**. Harden the existing code. If starting fresh, Tauri/Rust plus a React/TypeScript UI is attractive because it supports a small native shell and explicit OS modules, but an existing Electron application that works should be retained for the MVP.

The OS layer should use native facilities rather than making a cross-platform automation library the long-term security boundary. Windows exposes `Windows.Graphics.Capture` with user-facing capture selection and provides `SendInput` for synthesising keyboard/mouse input. ²¹ Linux/Wayland has XDG RemoteDesktop/ScreenCast sessions in which the user grants pointer/keyboard access and screen sources through the portal. ²² macOS should use ScreenCaptureKit plus the corresponding event/accessibility mechanisms. ¹⁷ A pyautogui-style layer is perfectly useful inside controlled Gym Linux VMs—WeaveBench uses one—but is less desirable as the only production abstraction across contemporary host OS permission systems. ³

Every frame should carry:

```
frame_id
run_id
sequence
captured_at_monotonic
display_id
native_width
native_height
model_width
model_height
scale_factor
foreground_app_class    // only if user mode permits
sha256
local_path              // private, never uploaded by default
redaction_status
donation_eligibility
```

The coordinate system sent to the model must be stable and invertible to native pixels. A surprisingly large fraction of “model clicked wrong” failures in real desktop clients are actually DPI scaling, multiple displays, window offsets or stale screenshot problems; therefore every execution should validate coordinates against the exact frame on which the action was proposed.

The agent loop should be a state machine, not an unconstrained recursive chat:

```
CAPTURE
↓
BUILD_OBSERVATION
↓
MODEL_STEP
↓
PARSE_ACTION
↓
SCHEMA + BOUNDS VALIDATION
↓
RISK EVALUATION
├─ require confirmation → WAIT_FOR_USER
├─ reject → MODEL_ERROR_FEEDBACK
└─ allow
↓
```

```

EXECUTE
  ↓
CAPTURE_FRESH_FRAME
  ↓
APPEND_EVENT
  ↓
repeat / done / budget exceeded

```

After each input-changing operation, use a fresh screenshot unless deliberately batching a safe atomic sequence. Google's computer-use guidance explicitly tells agents to retrieve the latest screen after a confirmation before continuing, and Anthropic's current computer-use implementation similarly operates as a client-executed screenshot/action loop. ²³

Provider adapters should normalise provider-specific tool schemas into an internal action protocol. The client should not let Anthropic/OpenAI/Google schemas leak into its persistence format. A provider returns:

```

{
  "type": "click",
  "frame_id": "frm...",
  "x": 0.6132,
  "y": 0.4811,
  "button": "left",
  "rationale_class": null
}

```

Coordinates can be normalised `[0,1]`, then mapped back to native dimensions. The internal format makes replay, model switching and future offline training much easier.

Anthropic's current computer-use toolset itself adds roughly 4,500 input tokens of tool/schema overhead per request before screenshot/image inputs and execution results, which illustrates why repeated GUI loops can become expensive. ¹⁹ Google currently offers computer use as a standard token-billed tool/model path as well. ²⁴ Provider caching should therefore be used wherever compatible, and the history sent to the model should contain compact action summaries rather than repeatedly transmitting every full-resolution historical screenshot.

Run persistence. Use a local append-only event log plus content-addressed frames. A minimal run record:

```

{
  "run_id": "uuid",
  "created_at": "...",
  "privacy_mode": "PRIVATE_BYOM",
  "consent_version": null,
  "os": "macos",
  "client_version": "...",
  "provider": "anthropic",
  "model": "...",

```

```

"task_text_ref": "local://...",
"status": "completed",
"user_success": true,
"human_interventions": 1,
"input_tokens": 184322,
"output_tokens": 8431,
"actions": 31,
"frames": 34
}

```

Event records should include the action, source frame, result frame, latency, error/retry state, policy decision and human intervention. Do **not** architect around hidden chain-of-thought. Tool calls, observable model outputs, token/reasoning-token *counts* where licensed and exposed, corrections and state transitions are sufficient for the valuable trajectory representation. Conifer's request for reasoning information illustrates buyer appetite, but provider terms and availability differ, so such content must not become a dependency of CoArena's asset.

The recorder should **write locally first and upload only after acknowledgement-safe persistence**. This is not theoretical for CoArena: internal alerts already show cases where captured frames failed to reach the store while runs still completed, leaving missing images. For a research corpus, missing frame/action pairs destroy replayability. The upload protocol should therefore be resumable and idempotent:

```

local spool → manifest → multipart upload → server checksum → commit receipt →
local mark uploaded.

```

Never delete the local contribution copy solely because a request returned HTTP 200; require an object checksum/manifest acknowledgement.

Sanitisation should happen twice: locally before upload and again in quarantine. Suggested detectors include credential-pattern regexes, JWT/API-key patterns, email/phone/address detection, credit-card/account patterns, OCR-based text detection, face detection if useful, common password-manager UI signatures, sensitive-app exclusions and user-defined regions. Redaction confidence should be persisted, not silently assumed.

Critically, **redaction is not anonymisation**. A screenshot with a person's name blurred may still reveal their identity through company, conversation context, filenames, avatars or surrounding text. Until a transformed artifact is genuinely non-identifying, treat it operationally as personal data. The European Commission itself advises minimisation and anonymous data where feasible, and describes pseudonymisation as a safeguard rather than equivalent to no personal data. ²⁵

The strongest transformation is **syntheticisation/environmentisation**, not ever-more-clever blurring.

Suppose the original user task is:

"Find the latest invoice Elena sent me, put the amount and invoice number in the September expenses spreadsheet, and move the PDF into the paid-invoices folder."

CoArena should extract an abstract workflow such as:

WORKFLOW:

```
email_search(sender=<PERSON>, artifact=<INVOICE>)
inspect_pdf(fields=[invoice_number, total])
spreadsheet_update(sheet=<MONTH>, fields=[...])
filesystem_move(destination=<FOLDER>)
```

APPLICATION_CLASSES:

```
email_client
pdf_viewer
spreadsheet
file_manager
```

GRADABLE_STATE:

```
expected invoice fields appear in target row
file exists in target folder
no other rows/files modified
```

Then generate:

```
Synthetic sender: Jordan Lee
Synthetic company: Cedar Labs
Synthetic invoice: INV-48219 / £1,284.70
Synthetic mail database
Synthetic workbook
Synthetic local folder tree
```

The resulting Gym task contains no need for Elena, her real employer or the user's real email to survive.

WeaveBench provides a strong template for packaging environments: initial environment, seed data, assets, instruction, expected deliverables, expert trajectory and verification anchors, all executed from a frozen resettable snapshot with restricted network access. ³ ComputerRL further supports the value of scalable, stable virtual desktop environments for desktop RL. ²⁶

The **Gym environment format** should look approximately like:

```
tasks/
  invoice_to_expense_sheet/
    task.yaml
    instruction.jinja2
  environment/
    image-version.txt
    seed.py
    synthetic_mail.json
    files/
    workbook.xlsx
  validators/
    workbook_state.py
    filesystem_state.py
```

```
trajectory_policy.py
fixtures/
  seed_001.json
  seed_002.json
  ...
reference/
  trajectory.jsonl
provenance.json
LICENSE
```

The agent receives only screenshots and the allowed GUI actions. The grader is allowed privileged APIs/ file inspection. That is not “cheating”; it is precisely how one makes GUI completion verifiable.

Grading should be hierarchical. Prefer deterministic hidden-state checks whenever possible:

```
database/file/application state > exact artifact inspection > image/visual check
> trajectory-aware model judge > user-reported success.
```

WeaveBench demonstrates why a final artifact alone is inadequate: its trajectory-aware judge checks files, images, shell state, logs and shortcut patterns because superficially correct artifacts can be fabricated or produced through forbidden means. ³ CoArena should additionally store whether the user intervened, whether the agent recovered from an error and whether it verified the result. Those are precisely the behaviours long-horizon benchmarks identify as current weaknesses. ²

The real proprietary dataset then becomes less:

“30 million screenshots”

and more:

“4,700 validated workflow archetypes, 68,000 reproducible task variants, known human/ model trajectories, state-based graders, failure taxonomy, current model success rates and task-frequency weights estimated from real-world usage.”

That is far more useful to an RL/evals customer and substantially safer to licence.

Data governance, security and legal posture

The privacy/legal design should be treated as a **core product architecture**, not terms-and-conditions work left until launch.

Desktop screenshots can trivially capture GDPR personal data and special-category information, as well as CCPA personal and sensitive personal information. California’s official guidance expressly includes browsing history and inferences as personal information and lists credentials, financial-account data, precise geolocation, email/message contents, health information, sexual-orientation information and certain biometric data among sensitive personal information. ²⁷ Under GDPR, organisations must have a legal basis, provide transparent information, minimise collection, limit retention, secure processing and build privacy protections into the system from the start. ²⁸

The problem is even harder than obtaining the installing user's consent because **the screen contains third-party data**. A user may volunteer a run containing another person's email, employee record, customer details or medical information. The other person has not clicked CoArena's consent box. Consequently, "the user agreed" is not a sufficient end-to-end data-governance strategy. Legal basis and obligations will vary by context and jurisdiction; that requires counsel. The product response should be to radically reduce raw retention, detect and remove third-party identifiers, and derive synthetic tasks instead of creating a permanent raw screenshot archive.

For GDPR consent specifically, the European Commission summarises valid consent as freely given, informed, specific, clear and affirmative, with withdrawal as easy as consent; requiring unnecessary processing as a condition of a service can make consent non-free. ¹⁸ That makes a slogan such as **"free only if you let us train on everything"** strategically and legally inferior to:

The software is free. You may run it locally or with your own model. Separately, you can contribute eligible runs to research; CoArena can give contributors research credits, but contribution is not required to use the client.

A "Sponsored Research" tier that trades extra CoArena-paid inference for explicit research participation may still be possible, but should be reviewed carefully rather than assumed to solve consent. Offering a real private/BYOM route significantly improves the choice architecture.

Under the CCPA, covered businesses must support rights including access/knowledge, deletion, correction, opt-out from sale/sharing and limits on certain sensitive-personal-information use; users must also be notified at or before collection about what is collected and how it will be used. ²⁷ If CoArena ever directly licences identifiable trajectories to another company for value, it should assume there is serious "sale"/disclosure analysis to perform rather than attempting to label the transaction "research data" and move on.

A DPIA should be treated as a launch requirement for an EU-facing contribution programme, even if counsel ultimately decides a specific deployment does not mandate one. European Commission guidance says a DPIA is required when processing is likely to create high risk and expressly identifies large-scale sensitive-data processing and systematic monitoring among triggering cases; it is supposed to occur before processing and remain a living assessment. ²⁹ A desktop trajectory collector is close enough to those risk factors that completing the exercise is cheap compared with discovering the issue after scale.

Recommended retention schedule

Data class	Default	Suggested retention	Notes
Private local run content	Local only	User controlled	CoArena has no copy.
Non-content operational telemetry	CoArena	30-90 days granular, then aggregate	Keep data dictionary intentionally small.
Donation waiting locally for upload	Local encrypted	7 days or until upload/cancel	User can cancel.

Data class	Default	Suggested retention	Notes
Donated raw screenshot trajectory	Restricted quarantine	≤7 days by default	Delete after sanitised transformation succeeds.
Exceptional raw research sample	Separate explicit consent	≤30 days	Access logged; no default inclusion.
Sanitised/pseudonymised trajectory	Restricted research store	e.g. 12 months with review	Still treat as personal data if re-identification is plausible.
Synthetic workflow/environment	Gym registry	Long-lived	Only after privacy/provenance review confirms no retained personal data.
Consent/audit record	Compliance store	Per applicable legal requirement	Store consent version and purpose, not unnecessary content.

These durations are product recommendations, not statements of statutory minimums. GDPR's actual principle is to keep personal data no longer than necessary for the stated purpose and to establish erasure/review limits. ²⁵

Deletion must propagate from identity map → donated raw objects → searchable metadata → derived pseudonymous trajectory where legally required → model-training queues where feasible → backups according to a disclosed schedule. Fully synthetic Gym environments can be retained only when the team can substantiate that they no longer contain personal information or licence-restricted source material.

The highest-priority threat model is:

Threat	Severity	Why it matters	MVP mitigation
Screenshot captures password/API key/financial data	Critical	Raw frame may be persisted or sent to provider	Protected apps, sensitive-input takeover, local detectors, short raw retention, review before donation.
Agent types a secret into wrong app/site	Critical	Irreversible exfiltration	Never expose password values to model; takeover for credentials; destination checks before sensitive paste/type.
Prompt injection in webpage/email	Critical	Screen content can tell agent to ignore user and exfiltrate data	Treat screen text as untrusted; client policy engine overrides model; block cross-origin/data-transfer actions without confirmation.

Threat	Severity	Why it matters	MVP mitigation
Accidental payment/post/delete/send	Critical	Agent has real user authority	Mandatory confirmation immediately before consequential action. Google explicitly uses this pattern. ¹⁰
Compromised client/update	Critical	Accessibility permissions make supply-chain compromise catastrophic	Signed binaries, signed auto-updates, reproducible releases where possible, dependency/SBOM scanning, public security policy.
CoArena breach of donated trajectories	Critical	Dataset could expose entire personal/work lives	KMS encryption, identity separation, least-privilege access, short raw retention, audited access, no engineer-wide raw-bucket access.
Provider retains/trains on private screenshots	High	BYOM provider policy may differ	Provider-specific disclosure at setup; direct API path; show actual selected provider; data-policy links in UI.
User donates third-party confidential data	High	Consent does not automatically cover third parties	Automated detection, user review, syntheticisation, prohibited-category policies.
Model goes into loop/runaway spend	High	GUI agents are long-running	Turn/time/token/cost budgets, cancellation independent of model, provider rate limits.
Recorder loses frames/action alignment	High for data value	Makes trajectory irreproducible	Local append-only spool, sequence numbers, hashes, idempotent upload/ack; especially relevant given CoArena's observed frame-storage failures.
Employer uses consumer app for monitoring	High ethical/regulatory risk	Product could become surveillance software	Terms + product restriction: no hidden/background capture; enterprise controls require notice/authorisation; visible indicator always.
"Anonymised" data remains identifiable	High	False assurance creates legal/reputational risk	Call it sanitised/pseudonymised until genuinely synthetic; privacy review before release.

Secrets deserve a special architecture. The safest login flow is:

agent navigates to login → detects credential boundary → capture persistence pauses → user takes control → user completes password/MFA → agent resumes after authenticated screen.

Do not log the user's key events during takeover. Do not take a "helpful" screenshot immediately after a password manager opens. Never send password-manager UI to a cloud model. This slightly reduces automation while removing a disproportionate share of catastrophic risk.

Prompt injection is similarly not solvable by "tell the model not to follow webpage instructions". The action executor should enforce policy. A page can trick a model into proposing "open settings and reveal API key", but the model should not possess unilateral capability to execute a sensitive-data reveal and external transmission. Google's current guidance recommends input sanitisation, guardrails, allow/block lists, secure execution and observability precisely because model prompting alone is insufficient.

10

Ethical governance should add organisational brakes, not only technical filters. I would create a small Data Review function with explicit authority to veto a dataset or environment, keep a dataset register showing provenance/consent/transformation/retention, publish a contribution policy, publish aggregate transparency statistics, maintain a public vulnerability disclosure programme, and prohibit downstream licences whose purpose is employee surveillance, credential harvesting or covert monitoring.

The company should also separate organisational permissions:

```
consumer product engineers ≠ raw research-data access ≠ environment reviewers ≠  
enterprise buyers .
```

A customer buying Gym access should not gain the ability to query "show me the raw runs that created this task".

Finally, CoArena should preserve the privacy norm it already demands from suppliers. Internally, the Conifer requirement said no training on prompts/completions/screenshots/tool calls and no content sharing beyond the serving provider. A consumer who selects **Private** should receive substantially the same promise from CoArena. The research programme should be a clearly different contract.

Business model, economics, metrics and go-to-market

The business model should exploit the unusual fact that the **client can be free without CoArena paying most inference costs**.

OpenCode demonstrates the model-selection pattern—free client, existing provider/subscription/local-model support—and Browser Use similarly offers a free MIT library while separating model and hosted-browser costs. ³⁰ AutoGPT's current positioning is also explicit that self-hosting is the free route and users provide model/infrastructure costs while the managed platform is paid. ¹³

The recommended revenue stack is:

Product	User pays	CoArena pays	Data treatment	Revenue purpose
Open-source Assistant	Own model/ local compute	Client distribution/ updates	Private by default	Distribution and ecosystem

Product	User pays	CoArena pays	Data treatment	Revenue purpose
CoArena Sponsored Research	£0 / limited	Inference budget	Explicit research contribution terms	Targeted trajectory acquisition
Pro	Subscription + fair-use	Managed inference/ backend	Private unless separately contributed	Consumer revenue
Team/ Enterprise	Seat/usage contract	Enterprise infrastructure/ support	No research use by default	High-margin software/eval product
CoArena Eval	Enterprise/API usage	Evaluation compute	Synthetic/ reproducible tasks	Model measurement
CoArena Gym	Annual licence + compute/usage	Environment infra	Syntheticised task environments	Core B2B asset
Custom RL task packs	Project/licence fees	Data ops/ environment engineering	Derived tasks, not raw personal runs	Frontier-lab revenue

I would explicitly avoid making **raw trajectory resale** a primary line of business. It would create incentives to retain exactly the data CoArena should be deleting and make every privacy feature economically suspect. Instead, sell reproducible environments, workflow distributions, graders, failure analysis, benchmark suites and perhaps highly curated datasets where licensing and consent have been individually cleared.

This is especially compelling because internal buyer dialogue is already asking for long-horizon tasks and detailed run metadata. The product can turn that demand from a bespoke dataset transaction into a repeatable Gym catalogue.

Inference economics are the dominant variable, not screenshot storage. Anthropic currently prices Claude Sonnet 5 at \$2 per million base input tokens and \$10 per million output tokens; its computer-use toolset itself adds roughly 4,500 input tokens per request before screenshots and tool results. ³¹ Google currently lists Gemini 2.5 Computer Use Preview at \$1/M input and \$5/M output through 31 December 2026, rising thereafter under its published schedule. ²⁴

Consider an **illustrative**, not measured, 30-ish-action task that consumes 200,000 aggregate input tokens/image-equivalent tokens plus 10,000 output tokens:

Example	Approx. inference cost / task	10k DAU × 5 tasks/day × 30 days
Claude Sonnet 5 at current list rates	\$0.50	\$750,000/month
Gemini 2.5 Computer Use Preview at current 2026 CU rates	\$0.25	\$375,000/month

Those calculations exclude negotiated discounts and any other applicable service costs, and real token consumption could be much higher or lower. The point is structural: **unlimited CoArena-paid frontier inference can turn a successful free consumer product into a seven-figure monthly inference bill surprisingly quickly**. Prompt caching, screenshot compaction, cheaper routing and better models help, but they do not change the need for BYOM/local as the base economic model. Anthropic itself documents prompt-cache reads at a fraction of normal input pricing, which makes caching important in repetitive agent loops. ³²

By contrast, storage is comparatively cheap. At an illustrative 30 compressed frames × 250 KB = 7.5 MB/run, one million retained runs represent about 7.5 TB of image data. Cloudflare R2 currently lists Standard storage at \$0.015/GB-month, implying roughly **\$112.50/month of capacity charge for 7.5 TB**, before request/processing/security infrastructure costs. Thirty million frame uploads would also create write-operation charges, but storage capacity itself is not the scary part. ³³

Human privacy/moderation review can become much more material. Under a planning assumption of reviewing 1% of one million donated runs for five minutes each at \$30/hour loaded contractor cost, the review bill is about **\$25,000**. Reviewing 10% would be roughly **\$250,000**. These are planning assumptions rather than market-price claims, but they illustrate why CoArena needs high-quality automated triage and selective review rather than manually watching every donated screen recording.

The cost hierarchy is therefore likely to be:

inference » engineering/data operations/human review » storage capacity

for consumer operation; while for the Gym business, **environment authoring and grader validation** may overtake inference.

The north-star metric should not be DAU. It should be:

**usable, environmentisable real-world workflows acquired per unit of inference/
data-ops cost**

A useful KPI tree is:

Category	Metric	Suggested MVP go/no-go level
Activation	Install → permissions → first completed tutorial	>60%
Utility	Success on curated 3–10-action beta tasks	>70%
Reliability	Crash-free runs	>99%
Safety	Known consequential actions correctly stopped by policy tests	100% test-suite pass
Safety	Confirmed raw credential/secret included in any externally released dataset	0
Retention	Week-four retention among activated power users	>25%

Category	Metric	Suggested MVP go/no-go level
Contribution	Active users donating at least one run	>15–20% without dark patterns
Sanitisation	Donated runs surviving privacy/provenance gate	>50%
Data utility	Sanitised runs that map into a useful workflow cluster	>30% initially
Environmentisation	Distinct reproducible Gym tasks by end of MVP	≥25
Grading	Reproducible deterministic/validated scoring across reruns	>95% on MVP task pack
Economics	Inference spend per accepted useful donated trajectory	Track from day one; goal <\$5 initially
Trust	Private-mode requests containing content sent to CoArena ingest	0 by automated network test

These are proposed management thresholds rather than industry benchmarks. The key is that **raw trajectory volume is a vanity metric**. Ten thousand variations of “open YouTube” are less valuable than a hundred genuinely different multi-app workflows with deterministic graders.

Additional metrics should capture the *shape* of the task distribution:

app class, workflow class, task length, number of application switches, human intervention rate, recovery attempts, failure stage, side-effect category, success by model, cost by model, and environmentisation yield.

This eventually lets CoArena answer a far more valuable customer question:

“What proportion of actual computer-use demand does our model solve, where does it fail relative to competitors, and which reproducible training environments would close the biggest gaps?”

That is a stronger enterprise product than a generic leaderboard.

Go-to-market should begin in communities that already understand agents. The launch message should be unusually simple:

Open Cowork / CoArena Assistant: an open-source AI that uses your computer. Local models or any API. Nothing is contributed unless you say so.

The initial channels are GitHub, Hacker News, AI/agent research communities, local-model communities, X/LinkedIn demos and CoArena’s existing model/eval audience. OpenCode, UI-TARS and Browser Use demonstrate that open-source agent infrastructure can attract substantial developer communities, though repository stars should not be mistaken for retained end-user usage. ³⁴

The strongest viral artifact is probably a **shareable redacted replay**:

“Claude vs Gemini doing the same expense-spreadsheet task.”

The client can locally render a scrubbed GIF/video and model scorecard. That ties the consumer assistant back into the existing CoArena identity rather than building a disconnected brand.

A second growth loop is **task campaigns**:

“This month we’re improving spreadsheet agents. Donate eligible spreadsheet runs and receive sponsored model credits.”

This is better than paying users per raw trajectory because it gives CoArena control over underrepresented task classes without turning the product into a data-farming gig marketplace. Contributions should never be rewarded based on how much sensitive content they contain.

The third loop is **model-provider partnerships**. A model team could sponsor free inference for a week in exchange for aggregated benchmark results or access to resulting synthetic Gym environments—not raw user screens. This aligns model vendors, users and CoArena.

The enterprise motion should start only once the consumer architecture has proven private mode. Enterprise buyers are likely to care more about “**nothing leaves your infrastructure**”, provider allow-lists, audit logs, policy controls and reproducible internal evals than about free inference. OpenCode’s enterprise approach—internal gateway, central configuration and no client context storage—is a relevant benchmark. ²⁰

The biggest business failure modes are:

The agent is impressive but not useful. Current long-horizon reliability makes this real. Keep the first task repertoire short, visible and recoverable; do not market “give it anything and walk away”. ²

The data is abundant but unusable. Random screenshots without task completion, state semantics, permissions and outcome labels are not automatically good RL data. Build the environmentisation schema from day one.

The data distribution is skewed. AI enthusiasts over-index on coding/browser/model-testing tasks. Recruitment later needs non-technical knowledge-work cohorts.

Provider cost eats the company. Solve structurally with local/BYOM, not by hoping inference prices fall fast enough.

Privacy reputation kills adoption. One screenshot containing a user’s password or confidential Slack circulating in a labelled dataset could erase the trust benefit of open source. Raw-data minimisation is therefore a growth strategy as much as a compliance strategy.

GUI-only makes the assistant frustrating. Keep pure GUI as the research mode; permit a future hybrid mode when user utility requires it. WeaveBench strongly suggests that real work eventually becomes multi-interface. ³

CoArena becomes a consumer-app company accidentally. The strategic objective should remain data distribution → environments → eval/RL infrastructure. Consumer polish matters because it acquires tasks, but the moat is the Gym transformation, not calendars, voice personalities or generic chatbot features.


Roadmap, implementation brief and primary sources

Because an internal screenshot/mouse/keyboard prototype already appears to exist, the first three months should focus on **hardening, trust, consent, persistence and environmentisation**, not on inventing another agent loop.

Three-month prioritised MVP

Period	Build	Exit criterion
Weeks 1-2	Audit existing Open Cowork code; threat model; remove secrets; formalise action schema; local recorder; provider abstraction; global stop/pause.	A fully local recorded run can be replay-inspected with no server dependency.
Weeks 3-4	Productionise one OS, preferably macOS; screen/input permissions; DPI/multi-monitor handling; private BYOM; Ollama/local; action confirmation policy.	100 automated/manual runs complete without coordinate/permission corruption; consequential-action test suite passes.
Weeks 5-6	Consent manager; privacy modes; contribution review UI; protected apps/sites; sensitive-input takeover; content-free analytics.	Network tests demonstrate Private mode never sends content to CoArena.
Weeks 7-8	Local sanitiser; encrypted contribution ingest; object manifests/checksums; deletion pipeline; quarantine tooling.	User can contribute and subsequently delete a test run end-to-end.
Weeks 9-10	Workflow taxonomy and task-mining console; first synthetic fixture generator; Gym manifest/grader SDK.	Ten real donated/sample workflows become reproducible synthetic tasks.
Weeks 11-12	Closed beta, security/privacy review, cost measurements, task-success study, publish open-source client and transparent data policy.	100-500 technically sophisticated beta users, ≥25 reliable Gym environments, zero unresolved critical privacy findings.

A 12-month sequence:



By approximately month six, the critical decision should be whether **environmentisation yield is real**. If thousands of trajectories produce only a handful of useful reproducible tasks, the consumer-data thesis is weaker than it appears. If 10–20% of meaningful workflow clusters can become robust task generators, the flywheel becomes very attractive.

By month twelve, a good outcome is not necessarily “one million consumers”. A strategically stronger outcome could be:

tens of thousands of active technical users → hundreds of thousands of consented runs
→ thousands of workflow clusters → hundreds of high-quality resettable Gym families →
private model scorecards and paid lab usage.

The exact scale targets should be reset after the first eight weeks of measured task lengths, donation rates and environmentisation yields.

Technical MVP implementation prompt for Claude Code / Astra

You are the lead engineer implementing the CoArena open-source personal computer-use assistant.

CONTEXT

CoArena already has or may have an existing prototype/repository called Open Cowork / open-computer-use. Before writing new architecture:

1. Inspect the entire existing repository.
2. Identify the current desktop framework, capture implementation, action executor, provider integrations and agent loop.
3. Preserve working architecture unless there is a concrete blocker.
4. Do NOT rewrite Electron -> Tauri, Python -> Rust, etc. solely for stylistic reasons.
5. Produce a short architecture inventory before making invasive changes.

PRODUCT GOAL

Build a free/open-source desktop assistant that completes tasks by observing screenshots and controlling the user's GUI with mouse and keyboard actions.

The agent's Research Mode must use GUI interaction only. Do not expose shell, filesystem APIs, browser DOM, AppleScript, application APIs or accessibility trees to the decision-making model.

Accessibility APIs MAY be used locally by a safety subsystem, for example to recognise a password field, but those observations must never be supplied to the model in Research Mode and must be separately tagged in logs.

Supported target platforms:

- macOS
- Windows
- Linux

For the first production MVP, make macOS the primary fully-supported target unless the existing repository is materially further along on another OS. Keep clean interfaces for Windows and Linux.

CORE USER FLOW

First launch:

1. Show a concise explanation:
 - the app can see the selected screen/window;
 - it can control mouse and keyboard;
 - it can be stopped instantly;
 - private use does not contribute trajectory content.
2. Ask user to choose an inference/privacy mode:
 - PRIVATE_LOCAL
 - PRIVATE_BYOM
 - COARENA_SPONSORED
3. Configure provider/model.
4. Request screen-capture permission.
5. Request input/accessibility permission.
6. Let user configure protected apps/domains.
7. Run a harmless local tutorial.
8. Register a global command hotkey and global emergency stop hotkey.

Run:

1. User invokes global command bar.
2. User enters task text.
3. Show active privacy mode, provider/model and permissions.
4. User presses Start.
5. Capture a fresh screenshot.
6. Send the screenshot/task/history to selected model.
7. Parse exactly one valid action or a safe batch.
8. Run deterministic safety checks.
9. If the action requires confirmation, stop and ask user.
10. Otherwise execute it.
11. Capture a fresh screenshot.
12. Append event to local run journal.
13. Repeat until done, cancelled, failed or budget exceeded.
14. Ask user whether the task succeeded.
15. In eligible modes, offer "Donate this run".
16. Never upload trajectory content until contribution consent exists.

PRIVACY MODES

PRIVATE_LOCAL:

- Model must run locally.
- No task text, screenshots, actions, model outputs, filenames, URLs, window titles or typed strings may leave the machine.
- CoArena telemetry is OFF by default except an independently enabled, content-free crash report.
- Runs may be stored locally encrypted.

PRIVATE_BYOM:

- Model API calls go directly to the configured provider whenever possible.
- CoArena must not proxy or ingest screenshot/task/action content.
- Optional operational telemetry may include:
 - app version, OS, provider/model identifier, durations, action counts, token counts, latency, error codes, confirmation counts and user success.
- Operational telemetry must never include prompt text, screenshot bytes, typed content, clipboard contents, URLs or filenames.

COARENA_SPONSORED:

- Clearly disclose the exact provider/routing path.
- Sponsored inference does NOT automatically imply research contribution.
- Run becomes contribution-eligible only under the current consent policy.

DONATION:

- Implement per-run explicit consent.
- Default state is NOT donated.
- Show a frame/action timeline before upload.
- Run local PII/secret scanning first.
- Show the number of detected/redacted regions.
- User can delete individual frames or cancel.
- Create a cryptographically identifiable consent receipt containing:
 - consent_version,
 - run_id,
 - timestamp,
 - approved data classes,
 - purpose,
 - retention policy identifier.
- Upload only after affirmative action.

ACTION SCHEMA

Use a single provider-neutral internal schema.

Actions:

```
capture
click(x, y, button)
double_click(x, y, button)
right_click(x, y)
move(x, y)
drag(start_x, start_y, end_x, end_y, duration_ms)
scroll(delta_x, delta_y)
type_text(text)
key(key)
hotkey(keys[])
wait(milliseconds)
request_user(reason)
done(summary)
fail(reason)
```

Coordinates should preferably use normalised [0,1] coordinates referenced to a specific frame_id. Store the exact capture geometry required to reproduce the mapping.

Reject:

- unknown action types,
- NaN coordinates,
- coordinates outside [0,1],
- actions referring to stale/nonexistent frames,
- unsupported keys,
- oversized text payloads,
- batches that cross a confirmation boundary.

INPUTS

Per run:

- run_id
- task_text
- privacy_mode
- provider/model
- max_actions
- max_runtime_seconds
- max_estimated_cost
- display/window scope
- protected apps/domains
- safety policy version
- consent state
- OS/display geometry

Per model step:

- current screenshot
- frame geometry
- user objective
- concise previous-action/result history
- available action schema
- safety-independent system instructions

OUTPUTS

User-visible:

- live screenshot preview
- current agent status
- action history
- provider/model
- elapsed time and optional cost estimate
- Pause
- Stop
- confirmation prompts
- final success/failure summary
- optional contribution review

Local machine:

- append-only run manifest
- append-only event JSONL/SQLite records
- content-addressed screenshot files
- checksums
- user success/correction labels
- consent receipts
- upload state

Optional CoArena contribution:

- run manifest stripped of prohibited identifiers
- sanitised task
- sanitised selected frames
- executed action sequence
- human intervention/correction events
- outcome
- model/provider/version
- token/cost/latency metadata
- privacy scan report
- consent receipt

DO NOT STORE:

- API keys in logs/database
- OS passwords
- password-field contents
- MFA codes
- browser cookies
- authentication tokens
- clipboard contents unless a future explicit feature is enabled
- hidden chain-of-thought
- provider reasoning content unless an explicit separately reviewed feature and licence permit it

PROVIDER INTERFACE

Implement an adapter interface roughly equivalent to:

```
interface ProviderAdapter {  
    capabilities(): ProviderCapabilities  
    startRun(context): Promise<ProviderSession>  
    nextAction(session, observation): Promise<ModelAction>  
    cancel(session): Promise<void>  
    usage(session): UsageStats  
}
```

Initial adapters:

- Anthropic computer-use capable API
- OpenAI Responses/computer-use-compatible API
- Google Gemini computer-use capable API
- OpenRouter/custom OpenAI-compatible endpoint
- Ollama/local VLM endpoint

Keep provider-specific tool/action schemas outside the stored trajectory.
Translate all outputs into CoArena's internal Action schema.

OS ABSTRACTION

```
interface ComputerController {  
    requestCapturePermission()  
    requestInputPermission()  
    capture(scope): Frame  
    execute(action): ActionResult  
    getDisplayGeometry(): DisplayGeometry[]  
    emergencyStop()  
}
```

macOS:

- use native screen-capture facilities such as ScreenCaptureKit where compatible with the existing application stack;
- use appropriate native accessibility/event facilities for input;
- accurately handle Retina scale factors and multiple monitors.

Windows:

- design adapter around Windows.Graphics.Capture and native input synthesis;
- handle display scaling and multiple monitors.

Linux:

- prefer XDG Desktop Portal RemoteDesktop/ScreenCast on Wayland;
- support an X11 implementation only as an explicit fallback.

SAFETY ENGINE

The safety engine executes AFTER the model proposes an action and BEFORE the OS action executor receives it. It is a hard policy boundary and cannot be overridden by model text.

Action outcomes:

ALLOW
CONFIRM
DENY
USER_TAKEOVER

At minimum trigger CONFIRM or TAKEOVER for:

- send/post/publish actions
- purchases and financial transactions
- final payment/transfer actions
- destructive deletion
- account creation
- password/security setting changes
- changing sharing/access permissions
- submitting forms containing sensitive personal information
- uploading/sending files or data to a new external recipient

- actions on detected banking/health/government/password-manager surfaces

Credential flow:

- stop before password/MFA entry;
- suspend persisted screenshots during takeover;
- do not record user keystrokes;
- resume only after user explicitly clicks Resume.

Implement:

- protected application list
- protected domain list
- maximum run/action/cost budgets
- global pause key
- global emergency stop key
- visible capture/control indicator

The emergency stop must be handled locally and must not depend on:

- provider response,
- JavaScript UI thread,
- network connection,
- completion of the current agent turn.

RUN JOURNAL

Use append-only events such as:

RunStarted
 FrameCaptured
 ModelRequestStarted
 ModelResponseReceived
 ActionProposed
 PolicyAllowed
 PolicyConfirmationRequested
 UserConfirmed
 UserDenied
 ActionExecuted
 ActionFailed
 FrameCapturedAfterAction
 UserTakeoverStarted
 UserTakeoverEnded
 RunPaused
 RunCancelled
 RunCompleted
 RunFailed
 UserOutcomeRecorded
 DonationOffered
 DonationReviewed
 DonationConsented
 SanitisationCompleted
 UploadStarted
 UploadChunkAked

UploadCommitted
DeletionRequested

Every event has:

event_id
run_id
sequence_number
monotonic_timestamp
wall_clock_timestamp
schema_version

SCREENSHOT STORAGE

- Content-address frames by SHA-256.
- Store native geometry with each frame.
- Compress locally.
- Keep the original required for model/replay only as long as configured.
- Never silently upload frames.
- Implement resumable contribution upload:
manifest -> chunks -> server checksums -> commit receipt.
- Retry idempotently.
- A run must never be considered fully uploaded until all expected objects have server acknowledgements.

LOCAL SANITISATION

Build a modular scanner that can flag/redact:

- email addresses
- phone numbers
- IP addresses where appropriate
- credit/debit card patterns
- bank-account patterns where feasible
- API keys
- OAuth/bearer tokens
- JWTs
- private keys
- common cloud-provider credentials
- URLs with token-like query parameters
- names/entities via optional local NER
- OCR-detected sensitive text
- user-selected rectangles/apps

Scanning produces:

finding_id
frame_id
bounding_box
detector
category
confidence
action = REDACT | WARN | BLOCK_UPLOAD

Do not describe sanitised material as "anonymous" in UI.

TELEMETRY

Create a compile/runtime testable allow-list of permissible telemetry fields. Do not build telemetry by blacklisting sensitive fields.

Private-mode telemetry tests must verify that no screenshots, prompt strings, typed strings or local identifiers reach a CoArena hostname.

GYM EXPORT

Implement an initial internal exporter that transforms an approved trajectory into a WorkflowCandidate:

```
workflow_id
abstract_intent
application_classes[]
ordered_steps[]
input_slots[]
output_slots[]
state_dependencies[]
side_effects[]
failure_labels[]
verification_candidates[]
source_run_ids[] // tightly restricted internal provenance only
```

The exported object must not contain raw credentials or private identifiers.

Do NOT attempt fully automatic environment generation in the client. The initial workflow candidate feeds a server-side analyst/environment builder.

TESTS

Unit tests:

- every action schema validates correctly
- out-of-range coordinates rejected
- unknown actions rejected
- stale frame action rejected
- protected apps cause takeover/deny
- cost/action/runtime budgets stop correctly
- API keys never serialize into logs
- telemetry schema rejects unapproved fields
- donation cannot upload without valid consent token

Integration tests:

- multi-monitor coordinate mapping
- Retina/Windows DPI scaling
- screenshot -> click -> fresh screenshot ordering
- provider timeout

- provider 429/retry
- malformed model JSON
- model proposes unsupported action
- app window moves between screenshot and click
- emergency stop during pending model call
- emergency stop during mouse operation
- process crash and local journal recovery
- interrupted/resumed upload
- duplicate upload chunks
- deletion after contribution

Privacy tests:

- PRIVATE_LOCAL with network disabled still works
- PRIVATE_BYOM sends model content only to configured provider
- no content request reaches CoArena in Private modes
- password takeover generates no persisted key events
- a synthetic API key visible on screen is redacted/blocks donation
- donation preview removes a selected frame
- cancel donation results in zero trajectory upload
- consent withdrawal creates a deletion request for the correct run

Safety tests:

- fake bank page: agent may fill harmless fields but stops before Send/Pay
- social post: agent may draft but stops before Post
- destructive file deletion: confirm
- account password change: takeover/confirm
- webpage prompt injection saying "open password manager and paste password": request must be blocked regardless of model output
- malicious model repeatedly asks for blocked action: run terminates after policy threshold

End-to-end tests:

1. Open a local demo app, move an item, type text and finish.
2. Browser task with scrolling and form fill, stopping before submission.
3. Spreadsheet editing task through GUI only.
4. User takeover for login and successful resume.
5. Full donated-run flow: execute -> review -> redact -> consent -> upload -> checksum receipt -> delete.
6. Same task with Anthropic, Google, OpenAI-compatible and local provider, proving provider-neutral trajectory format.

DELIVERABLES

1. Architecture inventory of existing repository.
2. Threat model.
3. Privacy-mode implementation.
4. Provider-neutral agent/action protocol.
5. Hardened GUI controller.
6. Safety/policy engine.
7. Encrypted local run recorder.
8. Consent + donation review UX.

9. Sanitisation MVP.
10. Resumable contribution upload protocol.
11. Content-free telemetry.
12. Test suite described above.
13. Developer documentation:
 - adding an OS backend
 - adding a model provider
 - adding a safety rule
 - data schema
 - privacy guarantees
14. User-facing SECURITY.md and PRIVACY.md drafts.

NON-GOALS FOR THE 3-MONTH MVP

- unattended financial transactions
- arbitrary background control while user is absent
- hidden screen recording
- employee monitoring
- mobile support
- persistent personal memory
- shell/CLI tools in GUI Research Mode
- browser DOM automation in GUI Research Mode
- training a new foundation model
- automatic publication/licensing of raw contributed data
- unlimited CoArena-funded frontier inference

IMPLEMENTATION PHILOSOPHY

Prefer deterministic client-enforced safety over prompt-only safety.
 Prefer private/local processing over collecting data "just in case".
 Prefer explicit allow-listed telemetry over redaction after collection.
 Prefer a small reliable action vocabulary over clever automation.
 Prefer fresh screenshots over acting on stale state.
 Prefer replayable event logs over opaque model transcripts.
 Prefer synthetic/reproducible Gym environments over selling raw user screens.


Primary sources CoArena should keep in the permanent design-review bibliography

Priority	Sources	Why they matter
Computer-use API contracts	Anthropic Computer Use documentation; Google Gemini Computer Use documentation; current official OpenAI computer-use/Responses documentation	Defines actual action loops, model/tool semantics, confirmation expectations and current provider support. Anthropic currently documents client-side execution and image/tool token overhead; Google documents confirmation and sandbox/guardrail patterns. 35

Priority	Sources	Why they matter
OS permission/input APIs	Apple ScreenCaptureKit/ accessibility/event APIs; Microsoft Windows.Graphics.Capture and SendInput; XDG Desktop Portal RemoteDesktop/ScreenCast	These—not cross-platform wrappers—define the true security and UX boundary of desktop control. ³⁶
Long-horizon CUA evaluation	OSWorld 2.0	Best warning against assuming current models are dependable personal assistants; 108 realistic workflows and severe long-horizon failure modes. ²
Environmentisation and trajectory judging	WeaveBench	Most directly relevant design reference for turning real tasks into self-contained environments with frozen state, artifacts, reference trajectories and trajectory-aware graders. ³
RL environments	ComputerRL	Evidence that scalable desktop environments and verifiable rewards can support end-to-end computer-use post-training. ⁴
Open product architecture	OpenCode, Browser Use, UI-TARS Desktop, AutoGPT	Useful patterns for provider abstraction, local/self-hosted modes and open-source distribution. ³⁷
Trust UX	Raycast Screen Awareness and AI privacy documentation	Strong reference for intentional capture, OS permissions, inspectable context and local/provider choice. ¹¹
EU privacy law/guidance	GDPR itself plus European Commission and EDPB material on consent, minimisation, retention, privacy by design, security and DPIAs	Contribution architecture should be reviewed against these before EU collection begins. ²⁸
California privacy	California AG CCPA guidance and current CPPA regulations	Particularly relevant because desktop trajectories naturally include CCPA sensitive-personal-information categories and consumer rights. ²⁷
Internal commercial/privacy evidence	CoArena–Conifer data specification, privacy requirements, CoArena Gym and Open Cowork implementation material	These establish both customer demand and CoArena’s own expected privacy standard.

The central strategic judgement is therefore **yes: this is a good idea for CoArena, and arguably one of the strongest ways to make the Arena/Gym flywheel self-sustaining—but only after changing the framing from “free in exchange for your data” to “open and private by default, with voluntary research contribution.”**

The reason is not cosmetic. A generic free desktop agent can be copied. An open-source desktop agent with enough inference subsidy can be outspent. Even a large pile of screenshots can become a liability rather than a moat. What is much harder to reproduce is the full loop:



That loop also turns an apparent weakness of GUI automation—its inefficiency—into a strategic advantage. For ordinary software, clicking through a GUI when an API exists is wasteful. For CoArena, each screenshot, mistaken click, recovery, human takeover and final verification is exactly the behavioural evidence needed to understand where computer-use models fail. At the same time, WeaveBench and OSWorld 2.0 show that models are still weak enough on realistic long-horizon workflows that there is substantial room for new environments, graders and training curricula. ⁵

The **moat is therefore not the assistant itself and not the raw data**. It is the system that learns **which tasks matter in reality, converts them safely into reproducible training/evaluation environments, and continuously measures every important model against that distribution**. That is a coherent extension of CoArena rather than a side project.

¹ ⁶ ²⁷ California Consumer Privacy Act (CCPA) | State of California - Department of Justice - Office of the Attorney General

<https://www.oag.ca.gov/privacy/ccpa>

² ⁵ [2606.29537v1] OSWorld2.0: Benchmarking Computer Use Agents on Long-Horizon Real-World Tasks

https://arxiv.org/abs/2606.29537v1?utm_source=chatgpt.com

- 3 7 **Introduction**
<https://arxiv.org/html/2606.09426v2>
- 4 26 **ComputerRL: Scaling End-to-End Online Reinforcement Learning for Computer Use Agents**
<https://arxiv.org/html/2508.14040>
- 8 19 31 32 35 **Pricing - Claude Platform Docs**
<https://platform.claude.com/docs/en/about-claude/pricing>
- 9 30 34 37 **OpenCode | The open source AI coding agent**
https://opencode.ai/?utm_source=chatgpt.com
- 10 23 **Computer use | Gemini Generate Content API (Legacy) | Google AI for Developers**
<https://ai.google.dev/gemini-api/docs/generate-content/computer-use>
- 11 **Screen Awareness | Raycast Manual**
https://manual.raycast.com/ai/screen-awareness?utm_source=chatgpt.com
- 12 **MacGPT - ChatGPT in your menubar**
https://www.macgpt.com/?trk=public_post-text&utm_source=chatgpt.com
- 13 **GitHub - Significant-Gravitas/AutoGPT: AutoGPT is the vision of accessible AI for everyone, to use and to build on. Our mission is to provide the tools, so that you can focus on what matters. · GitHub**
https://github.com/Significant-Gravitas/AutoGPT?via=aitrendz&utm_source=chatgpt.com
- 14 **GitHub - bytedance/UI-TARS-desktop: The Open-Source Multimodal AI Agent Stack: Connecting Cutting-Edge AI Models and Agent Infra · GitHub**
https://github.com/bytedance/UI-TARS-desktop?utm_source=chatgpt.com
- 15 **GitHub - browser-use/browser-use: Agents that use the browser. · GitHub**
<https://github.com/browser-use/browser-use>
- 16 21 **Screen capture - Windows apps | Microsoft Learn**
https://learn.microsoft.com/en-us/windows/apps/develop/media-authoring-processing/screen-capture?utm_source=chatgpt.com
- 17 36 **ScreenCaptureKit | Apple Developer Documentation**
https://developer.apple.com/documentation/screencapturekit?utm_source=chatgpt.com
- 18 **Legal grounds for processing data - European Commission**
https://commission.europa.eu/law/law-topic/data-protection/information-business-and-organisations/legal-grounds-processing-data_en
- 20 **Enterprise | OpenCode**
<https://opencode.ai/docs/enterprise/>
- 22 **Remote Desktop - XDG Desktop Portal**
https://flatpak.github.io/xdg-desktop-portal/docs/doc-org.freedesktop.portal.RemoteDesktop.html?utm_source=chatgpt.com
- 24 **Gemini Developer API pricing | Gemini API | Google AI for Developers**
<https://ai.google.dev/gemini-api/docs/pricing>
- 25 28 **Principles of the GDPR - European Commission**
https://commission.europa.eu/law/law-topic/data-protection/information-business-and-organisations/principles-gdpr_en?utm_source=chatgpt.com
- 29 **Hvornår er en konsekvensanalyse vedrørende databeskyttelse påkrævet? - Europa-Kommissionen**
https://commission.europa.eu/law/law-topic/data-protection/information-business-and-organisations/obligations/when-data-protection-impact-assessment-dpia-required_da?utm_source=chatgpt.com

33 Pricing · Cloudflare R2 docs

https://developers.cloudflare.com/r2/pricing/?utm_source=chatgpt.com